

# sparklyR: R and Spark

## An Introduction

Marck Vaisman

November 26, 2016

# Disclaimer

*The opinions expressed in this presentation are mine and they do not reflect in any way those of the organizations which I am affiliated with.*

## Show of hands

- ▶ Who has experience with Hadoop?

## Show of hands

- ▶ Who has experience with Hadoop?
- ▶ Who has experience with Spark?

## Show of hands

- ▶ Who has experience with Hadoop?
- ▶ Who has experience with Spark?
- ▶ Who has experience with Scala?

# Show of hands

- ▶ Who has experience with Hadoop?
- ▶ Who has experience with Spark?
- ▶ Who has experience with Scala?
- ▶ Who **loves** R?

# Agenda

- ▶ About Spark
- ▶ Introduce sparklyr and compare it to SparkR
- ▶ Show how to install and configure so everything works together nicely on AWS EMR
- ▶ Quick Demo

# Spark

- ▶ An Apache project providing **lightning fast cluster computing**
- ▶ Open source cluster computing framework
- ▶ Works with data in memory as opposed to batch io from disk
- ▶ Developed in Scala
- ▶ Provides APIs in Scala, Java, Python and R
- ▶ Can run on clusters managed by YARN, Mesos or Standalone
- ▶ Lazy evaluation: no computations are performed until an action is taken or the data is collected back to the driver (master)
- ▶ Awesome, but finicky



## sparklyr

- ▶ A new package developed by the RStudio team (enough said)
- ▶ Provides a complete dplyr interface to Spark RDDs
- ▶ Transforms dplyr verbs into SparkSQL commands run that act on a Spark DataFrame
- ▶ **Not** a replacement to SparkR
- ▶ Lowers barrier to entry into Spark for R users

## Easy Button!

From this (native Scala on spark-shell)

```
val queries = sc.textFile("hdfs://myfile")
val tups = queries.map(line => line.split('\t'))
val countHour = tups.map(x => (x(0) + "+" + x(2).take(13) + "+"))
val byUserHour = countHour.map(x => (x._1.split("\\+")(0), x._1.split("\\+")(1)))
val byUser = byUserHour.groupByKey // RDD[(String, Iterable[(String, String)])]
val times = countHour.map(x => x._1.split("\\+")(1)).distinct()
val broadcastTimes = sc.broadcast(times)
```

To this

```
library(sparklyr)
library(dplyr)
sc <- spark_connect(master = "yarn-client")
queries <- spark_read_csv(sc = sc, name = "green",
                           path = "hdfs://myfile")
```

# sparklyr Functions

- ▶ `spark_` Connecting to Spark
- ▶ `spark_read_` and `spark_write_` Read and write Spark DataFrames from CSV, JSON and Parquet
- ▶ `sdf_` Operations on Spark DataFrames
- ▶ `ml_` Functions to invoke ML algos
- ▶ `ft_` Functions to transform Spark DataFrames
- ▶ Documentation

# sparklyR vs. SparkR from Stack Overflow

## sparklyr

- ▶ Does not support `do()` - arbitrary functions on groups or rows

## SparkR (built into Spark 1.6+)

- ▶ More general front end to Spark using R

# Installing sparklyr (well, other things too)

## Use a Bootstrap Action

- ▶ Installs RStudio Server on Master node of EMR cluster
- ▶ Accessible via Web at `ec2-ip-of-master-node.amazon.com:8787`
- ▶ Login as hadoop user (with hadoop password)
- ▶ Sets all environment variables
- ▶ Also installs sparklyr, SparkR and other options

## Caveats of this BA

- ▶ Older version of RStudio
- ▶ Specific to AWS EMR Hadoop configuration

# Using sparklyr

## Cluster

```
# This assumes your environment variables are already set  
library(sparklyr)  
library(dplyr)  
  
# To connect to a cluster  
sc <- spark_connect(master = "yarn-client")
```

## Local

```
Sys.setenv(SPARK_HOME="/usr/local/Cellar/apache-spark/2.0.2")  
library(sparklyr)  
library(dplyr)  
  
sc <- spark_connect(master = "local")
```

# Analyzing NYC Taxi Data (the new Big Data iris dataset)

Github Repo from RStudio Example

## Setup

- ▶ Downloaded the NYC Taxi dataset used by (Todd Schneider)[<http://toddschneider.com/posts/analyzing-1-1-billion-nyc-taxi-and-uber-trips-with-a-v>] and put it on s3 (s3://bigdatateaching)
- ▶ Created Parquet files from the CSVs using sparklyr

## What we'll see

- ▶ Read in Parquet files and create Spark DataFrames
- ▶ Do some simple munging and aggregations

## Spark Parameters to tweak (YMMV)

- ▶ BA Used allocates all available resources to Spark
- ▶ `spark.driver.cores`
- ▶ `spark.executor.cores`
- ▶ `spark.executor.memory`



## Additional Resources

- ▶ [<http://www.agildata.com/apache-spark-2-0-api-improvements-rdd-dataframe-dataset>]
- ▶ [<http://blog.revolutionanalytics.com/2016/10/tutorial-scalable-r-on-spark.html>]
- ▶ [<https://www.toptal.com/spark/introduction-to-apache-spark>]
- ▶ [<https://0x0fff.com/spark-architecture/>]
- ▶ [<https://0x0fff.com/spark-architecture-shuffle/>]
- ▶ [<https://0x0fff.com/spark-memory-management/>]
- ▶ [<https://0x0fff.com/apache-spark-future/>]

## Parting thoughts

- ▶ Like all Big Data tools, Spark is finicky and you should understand more about the internal workings of Spark. It is very powerful and fast when configured and used correctly
- ▶ I've been collecting large datasets that are publicly available but not necessarily readily available and/or usable on Amazon. Hosted on `s3://bigdatateaching/`
  - ▶ NYC Taxi Data (csv files and Parquet)
  - ▶ Criteo 1TB Dataset
  - ▶ More to come
- ▶ Also working on making the *ideal Data Science environment* setup easier using Ansible, which can be used on many platforms

# Thank you!

- ▶ `vaisman_marck@bah.com`
- ▶ `marck@datacommunitydc.org`
- ▶ `twitter.com/wahalulu`